

In [1]:

```
#importing libraries  
import numpy as np  
import matplotlib.pyplot as plt  
import random  
from tqdm import tqdm
```

In [2]:

```

#building the environment
class Environment:
    def __init__(self, M,N, holes, terminal_state): #initializing the environment state
s, holes, terminals and rewards
        self.states = set()
        self.shape = (M,N)
        self.holes = holes
        self.terminal_state = terminal_state

        for i in range(1,M+1):
            for j in range(1, N+1):
                if (i,j) not in holes:
                    self.states.add((i,j))

        self.rewards = self.initialize_rewards()
        self.prob_agent_action = [0.8, 0.1, 0.05, 0.05]
    def initialize_rewards(self): #function to initialize the rewards for each state of
the environment
        r = {}
        for state in self.states:
            if state == (6,3):
                r[state] = -15
            elif state == (6,6):
                r[state] = 15
            else:
                r[state] = 0

        return r

    def agent_move(self, s, a): #function to update the state of the agent given an act
ion a and current state s
        x, y = s
        if a=='U':
            x = x-1
        elif a=='D':
            x = x + 1
        elif a=='R':
            y = y + 1
        elif a=='L':
            y = y - 1

        stay_same = self.check_corner_and_hole((x,y))

        if stay_same:
            return s

        return (x,y)

    def move_clockwise90(self, a): #function to return the action which is a 90 degree
rotation to current action a
        if a=='U':
            return 'R'
        elif a=='R':
            return 'D'
        elif a=='D':
            return 'L'
        elif a=='L':
            return 'U'

```

```
def move_anti_clockwise90(self, a): #function to return the action which is a 90 de  
gree rotation to current action a  
    if a=='U':  
        return 'L'  
    elif a=='L':  
        return 'D'  
    elif a=='D':  
        return 'R'  
    elif a=='R':  
        return 'U'  
  
def check_corner_and_hole(self, s):  
    #function to check if the updates state goes out of the gridworld or goes into  
holes.  
    #If so, it returns a True value to address that the update should not take plac  
e and agent should remain in current state.  
    x1, y1 = s  
    stay_same = False  
    for hole in self.holes:  
        if (x1,y1) == hole:  
            stay_same = True  
  
    if x1<1 or x1>6:  
        stay_same = True  
    if y1<1 or y1>6:  
        stay_same = True  
  
    return stay_same
```

In [3]:

```
gridworld = Environment(6,6, [(4,3),(5,3)], (6,6))
```

In [7]:

```

#building the cross entropy method agent
class CEM:
    #initializing the agent parameters
    def __init__(self, gamma, env , K0 = 500, epsilon = 0.25):
        self.actions = ['L','R','U','D'] #possible actions
        self.gamma = gamma #discount parameter
        self.K = K0
        self.param_dim = len(env.states) * len(self.actions)
        self.param_mean = [10] * self.param_dim
        self.param_cov = 5*np.identity(self.param_dim)

        self.epsilon = epsilon
        self.initialize_theta(env)
#        self.V = self.initialize_value_states(env)
#        self.Q = self.initialize_Qvalue(env)

    def initialize_theta(self,env):
        #function to intialize K thetas where each theta is sampled from a multivariate
        gaussian
        self.state2param = {}
        self.param2state = {}
        i = 0
        for s in env.states:
            for a in self.actions:
                self.state2param[(s,a)] = i
                self.param2state[i] = (s,a)
                i+=1

        self.thetas = np.random.multivariate_normal(self.param_mean, self.param_cov, si
        ze = self.K)
        self.pi = self.find_policy(self.param_mean, env)
        #print (self.thetas.shape)

    def possible_actions(self, s, pi):
        #function to return the possible actions given a state s and policy pi
        actions = [a for a in pi[s]]
        return actions

    def initialize_Qvalue(self, env):
        #function to initialize Q values
        Q = {}
        for s in env.states:
            Q[s] = {}
            for a in self.actions:
                Q[s][a] = 0

        return Q

    def initialize_value_states(self, env):
        #function to inititalize value of states
        v_s = {}
        for state in env.states:
            v_s[state] = 0

        return v_s

    def find_policy(self,theta1,env):
        #function to find the policy corresponding to a given theta which is a vector o
        f parameters of size 136 (=34*4)

```

```

pi = {}
for s in env.states:
    pi[s] = {}

for i in range(self.param_dim):
    s,a = self.param2state[i]
    theta = theta1[i]
    pi[s][a] = theta

```

```

for s in env.states:
    prob_sum = 0
    for a in self.actions:
        p = np.exp(pi[s][a])
        prob_sum+=p
        pi[s][a] = p

    for a in self.actions:
        pi[s][a] /= prob_sum

```

```

return pi

```

```

def compute_value(self, env, pi, epsilon = 1e-8):
    #function to evaluate the policy pi
    delta = np.inf
    Q = {}
    V = {}
    for s in env.states:
        Q[s] = {}
        if s==env.terminal_state:
            V[s] = 15
        else:
            V[s] = random.random()

        for a in self.actions:
            Q[s][a] = 0

    while delta > epsilon:
        max_diff = -np.inf
        for s in env.states:
            if s!=env.terminal_state:
                v = 0
                for a in pi[s]:
                    p = pi[s][a]
                    #with prob 0.8 take action a to get to state s1 with reward r1
                    s1 = env.agent_move(s,a)
                    r1 = env.rewards[s1]

                    v1 = 0.8 * (r1 + self.gamma*V[s1])

                    #with prob 0.1 stay in same state s
                    r2 = env.rewards[s]
                    v2 = 0.1 * (r2 + self.gamma * V[s])

                    #with prob 0.05 take action in direction +90 degree clockwise d
                    irection
                    a3 = env.move_clockwise90(a)
                    s3 = env.agent_move(s, a3)

```

```

r3 = env.rewards[s3]
v3 = 0.05 * (r3 + self.gamma * V[s3])

#with prob 0.05 take action in direction -90 degree clockwise d
irection

a4 = env.move_anti_clockwise90(a)
s4 = env.agent_move(s, a4)
r4 = env.rewards[s4]
v4 = 0.05 * (r4 + self.gamma * V[s4])

v = v + (p * (v1+v2+v3+v4))
Q[s][a] = v1+v2+v3+v4

diff = abs(V[s] - v)
V[s] = v
max_diff = max(max_diff, diff)

delta = max_diff

return V,Q

def param_update(self, env, keep_pct = 0.9):
    #function to update parameters of the cross entropy method

    #the new K value will be maximum of 90% of current K value and 80
    K_new = max(np.int(keep_pct*self.K),80)
    values = []
    max_s11_value = -np.inf

    #for each theta, find policy corresponding to it and evaluate the Q values
    Qs = []
    for i in range(self.K):
        theta_i = self.thetas[i,:]
        pi = self.find_policy(theta_i, env)
        V,Q = self.compute_value(env, pi)
        Qs.append(Q)
        s11_value = V[(1,1)]
        #tot_val = self.compute_total_state_value(V, env)

        if s11_value > max_s11_value:
            max_s11_value = s11_value
            self.best_theta = theta_i

    values.append(s11_value)

    # best_indices = np.argsort(values)[-K_new:]
    # thetas = self.thetas[best_indices,]

    #for each parameter of theta corressponding to a state action pair choose K_new
    theta values which has the best Q values
    thetas = []
    for i in range(self.param_dim):
        s,a = self.param2state[i]
        qsa_values = [q[s][a] for q in Qs]
        best_indices = np.argsort(qsa_values)[-K_new:]
        theta_sa_new = list(self.thetas[best_indices,i])
        thetas.append(theta_sa_new)

    thetas = np.array(thetas).T

```

```

self.param_mean, self.param_cov = self.update_mean_and_cov(thetas)

#Update mean and covariance vector of theta based on these new K values
self.param_mean = np.mean(thetas, axis = 0)
self.param_cov = np.cov(thetas.T)

self.K = K_new

# Sample a new K theta values based on updated parameters
self.thetas = np.random.multivariate_normal(self.param_mean, self.param_cov, si
ze = self.K)
self.pi = self.find_policy(self.param_mean, env)

#print (max_s11_value)
return max_s11_value

def play(self, env, iterations = 100):
    #function to play the cross entropy method over 100 iterations
    v_s11 = []
    for iteration in tqdm(range(iterations)):
        s11 = self.param_update(env)
        if iteration==(iterations-1):
            self.pi = self.find_policy(self.best_theta, env)

        v_s11.append(s11)

    return v_s11

def find_optimal_move(self, s):
    #function to find the optimal move given the policy
    optimal_moves = []
    max_v = 0
    for a in self.pi[s]:
        if self.pi[s][a] > max_v:
            optimal_moves = [a]
            max_v = self.pi[s][a]
        elif self.pi[s][a] == max_v:
            optimal_moves.append(a)

    return optimal_moves

def plot_state_and_policy(self, env, plot_state_values = False, label = None):
    #function to plot state values and policy
    plt.figure(figsize = (20, 10))
    Grid_plot=plt.subplot()
    M,N = env.shape
    for i in range(M):
        for j in range(N):
            s = (i+1, j+1)
            if s==env.terminal_state:
                if plot_state_values:
                    t = round(self.V[s], 3)
                    value = str(s) + "\n\n" + str(t) + "\n\n" + "TERMINAL"
                else:
                    value = "TERMINAL"
            elif s not in env.holes:
                value=str(s)
                move = self.find_optimal_move(s)
                if plot_state_values==False:
                    value = value + '\n\n' + ', '.join(move)

```

```
Grid_plot.text(j+0.5,N-i-0.5,value,ha='center',va='center')
```

```
agent = CEM(0.9, gridworld)
```

```
values = agent.play(gridworld)
```

```
100%|███████████████████████████████████████████████████████████████████████████|  
██████████ | 100/100 [17:37<00:00, 10.57s/it]
```

### Plot of values of state (1,1) across iterations of Cross entropy method



In [10]:

```
def plot_values(values, label = None):  
    ts = np.arange(1, (len(values)+1))  
    plt.plot(ts, values)  
    plt.xlabel("Iteration")  
    plt.ylabel("Value of state (1,1)")  
    plt.title("V[(1,1)] across iterations")  
    if label is not None:  
        plt.savefig(label)  
    plt.show()  
  
plot_values(values, label = 'plots/1_c_value_state_across_iterations.png')
```



Plot of policy learnt by the cross entropy method

In [12]:

```
agent.plot_state_and_policy(gridworld, label = 'plots/1_c_optimal_policy.png')
```

(1, 1) R	(1, 2) R	(1, 3) D	(1, 4) R	(1, 5) D	(1, 6) D
(2, 1) R	(2, 2) D	(2, 3) R	(2, 4) D	(2, 5) D	(2, 6) D
(3, 1) R	(3, 2) R	(3, 3) U	(3, 4) R	(3, 5) D	(3, 6) D
(4, 1) U	(4, 2) U	HOLE	(4, 4) R	(4, 5) R	(4, 6) D
(5, 1) R	(5, 2) U	HOLE	(5, 4) R	(5, 5) R	(5, 6) D
(6, 1) L	(6, 2) L	(6, 3) R	(6, 4) R	(6, 5) R	TERMINAL

In [ ]: